

PROJECT

Identify Fraud from Enron Email

A part of the Data Analyst Nanodegree Program

PROJECT REVIEW

CODE REVIEW

NOTES

SHARE YOUR ACCOMPLISHMENT!  

Meets Specifications

IMPORTANT: You were absolutely right on your Chi-Squared test doubts! I was wrong assuming email ratio features were discrete features also, sorry for that 😞. So, a proper test for this case would be a t-test, the main goal of the test is to test if two groups are similar or not, in this case, groups would be the email ratio distributions for non_POIs and POIs. This is a [great post](#) describing a step-by-step t-test implementation in python!.

Hope this resolves your question, but if not, please email us and we can continue discussing this topic: dataanalyst-support@udacity.com

Great work!. Yours is a strong project!. You demonstrate a good understanding of Machine Learning and your report is written in explanatory terms allowing your audience to follow and understand the work done.

Congratulations on passing your exam and stay Udacious!

Quality of Code

Code reflects the description in the answers to questions in the writeup. i.e. code performs the functions documented in the writeup and the writeup clearly specifies the final analysis strategy.

Your written response perfectly describes your strategy, includes the required level of detail when required and your code includes all the processes mentioned. Machine learning is not just about building good models, it is also about communicating results to your audience in a clear and direct way. You achieve both goals. Well done!.

As a suggestion, this project is a great opportunity for you to create a new repository in Github that becomes part of your online portfolio and allow potential employers to review your work. This report defines your credentials, so it is important that you put special attention not just to the technical side of the project but also the communications side since this is a critical characteristic for any data scientist. For your reference, check this [Kaggle post](#) for further reference, as you can see this is really a hot topic in the data science world! 😊

Since I reviewed your previous submission, I maintain my comments for the unchanged sections.

`poi_id.py` can be run to export the dataset, list of features and algorithm, so that the final algorithm can be checked easily using `tester.py`.

All required `.pkl` files are included and `poi_id.py` worked without problems.

Understanding the Dataset and Question

Student response addresses the most important characteristics of the dataset and uses these characteristics to inform their analysis. Important characteristics include:

- total number of data points
- allocation across classes (POI/non-POI)
- number of features used
- are there features with many missing values? etc.

Your report includes a description of the main characteristics of the dataset.

Note these numbers are particularly important since they describe the main dataset characteristics:

1. The small data set is why the tester.py file uses StratifiedShuffleSplit instead of a simpler cross-validation method such as TrainTestSplit. StratifiedShuffleSplit will make randomly chosen training and test sets multiple times and average the results over all the tests.
2. The data is unbalanced with many more non-POIs than POIs. StratifiedShuffleSplit also makes sure that the ratio of non-POI:POI is the same in the training and test sets as it was in the larger data set.
3. The unbalanced data is also why we use precision and recall instead of accuracy as our evaluation metric.

For your reference, some [techniques](#) to handle unbalanced datasets and this [repo](#) with plenty of tools.

Student response identifies outlier(s) in the financial data, and explains how they are removed or otherwise handled.

Nice job finding the TOTAL row. As a suggestion, the dataset comes in a json format that makes it difficult to handle and explore. My suggestion is to use [Pandas](#):

```
data_dict = pickle.load(open("final_project_dataset.pkl", "r") )
###creating dataframe from dictionary - pandas
df = pandas.DataFrame.from_dict(data_dict, orient='index', dtype=np.float)
print df.describe().loc[:, ['salary', 'bonus']]
```

Check the [docs](#) for more information on how pandas can read data from different sources.

Optimize Feature Selection/Engineering

At least one new feature is implemented. Justification for that feature is provided in the written response, and the effect of that feature on the final algorithm performance is tested. The student is not required to include their new feature in their final feature set.

Good work engineering your features, including your reasons and testing their impact over your classifier.

Univariate or recursive feature selection is deployed, or features are selected by hand (different combinations of features are attempted, and the performance is documented for each one). Features that are selected are reported and the number of features selected is justified. For an algorithm that supports getting the feature importances (e.g. decision tree) or feature scores (e.g. SelectKBest), those are documented as well.

Well done including the PCA components variances associated!

If algorithm calls for scaled features, feature scaling is deployed.

It is great you scaled your features since some of the classifiers attempted calls for it. However, it would be great if you could include few words in your written response describing which of these algorithms call for scaled features.

Also, great work using a log transform for your features and testing its impact! 👍

Pick and Tune an Algorithm

At least 2 different algorithms are attempted and their performance is compared, with the more performant one used in the final analysis.

Outstanding work testing several algorithms and comparing their performance in terms of f1 using different validation techniques.

As a side comment, note F1 is a subclass of the F-Scores as [here](#) described. For example, F0.5 puts more importance over Precision than Recall, it is up to your problems needs to decide any F-Score between 0 (only considers Precision) and infinite (only consider Recall). [Here](#) you can find more information.

Response addresses what it means to perform parameter tuning and why it is important.

Good understanding of tuning.

At least one important parameter tuned with at least 3 settings investigated systematically, or any of the following are true:

- GridSearchCV used for parameter tuning
- Several parameters tuned
- Parameter tuning incorporated into algorithm selection (i.e. parameters tuned for more than one algorithm, and best algorithm-tune combination selected for final analysis).

Excellent use of GridCV to tune your classifier, I especially like how you optimized it by passing the right CV object and oriented your search to maximize the score of your choice.

Also, wise decision using [RandomizedGridCV](#) is a good alternative (much faster than GridSearchCV) since it explores a subset of points in the parameter space defined (specially important when working with huge datasets).

For your reference, note you can also [visualize your grid results](#).

Validate and Evaluate

At least two appropriate metrics are used to evaluate algorithm performance (e.g. precision and recall), and the student articulates what those metrics measure in context of the project task.

Good work using precision&recall to evaluate your classifier. Also, good definition of both in terms of POI detection.

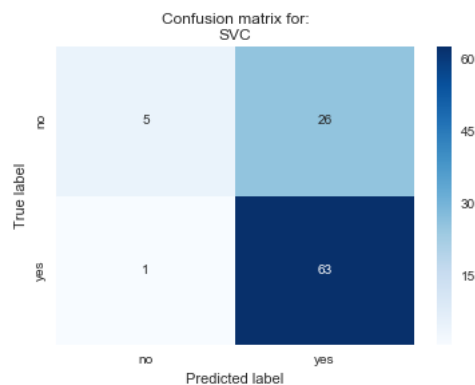
As a side comment: Note accuracy is not a good score in this case since the dataset in this project is very small and the ratio of negatives to positives is highly skewed (18 POIs out of 146 total cases), a classifier that predicted only non-POIs as output, would obtain an accuracy score of 87.4%. In this regard, accuracy is not a strong validation metric for our classifier. For your reference, [this interesting post](#).

As a suggestion, consider a [confusion matrix](#) for more detailed information about classifier performance.

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

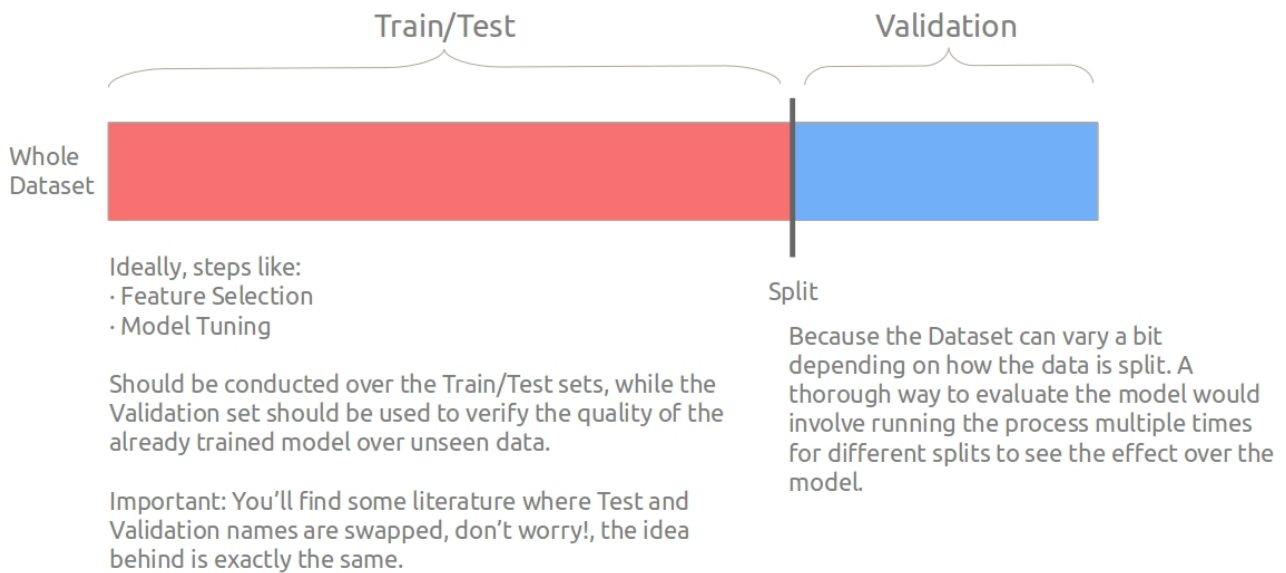
# Compute confusion matrix for a model
model = clf_C
cm = confusion_matrix(y_test.values, model.predict(X_test))

# view with a heatmap
sns.heatmap(cm, annot=True, cmap='Blues', xticklabels=['no', 'yes'], yticklabels=['no', 'yes'])
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.title('Confusion matrix for:\n{}'.format(model.__class__.__name__));
```



Response addresses what validation is and why it is important.

Good understanding of validation. As a side comment, since this dataset is really really small, validation is performed just using the test set. In the real world, you would split your dataset in train/test/validation in order to validate your trained&tuned model against completely unseen data. For further reference, check [this excellent answer](#).



Performance of the final algorithm selected is assessed by splitting the data into training and testing sets or through the use of cross validation, noting the specific type of validation performed.

For this particular problem, a stratified shuffle split is preferred for a number of reasons. Since the dataset is small, a shuffle split will randomly assign entries to test and training sets in a fold. However, the stratification preserves the percentage in the target class as in the complete set. This is particularly important Validation for our investigation as we have such a small percentage of Pols - if we did not stratify we could easily have a test or train set which contained no entries that were Pols. If we had no entries that were Pols in either a test or train set then the model would perform very badly.

When `tester.py` is used to evaluate performance, precision and recall are both at least 0.3.

Well done!. Your classifier is over 0.3 for precision&recall scores.

[↓ DOWNLOAD PROJECT](#)

Have a question about your review? Email us at review-support@udacity.com and include the link to this review.

RETURN TO PATH

Rate this review

[Student FAQ](#)